Capstone Checkpoint 5.1 - Agentic Retriever Loop

Oracle, vector, and hybrid experiment comparison. Generated 2026-09-21 10:31 EDT.

7	2	124	19	0
Included runs	Configurations	Question-runs	Questions	Legacy runs excluded

Summary

The highest observed mean quality score was 0.754 for vs_k7_md_graph_n5retriever across 3 run(s). Quality assigns 1 to pass, 0.5 to partial, and 0 to fail.

The highest observed mean designated-source recall was 93.9% for vs_k7_md_graph_n5retriever. Article recall is recomputed from the required and retrieved article names in every CSV.

The highest task-family attention score was 73.3 for Exhaustive list or time filter under vs_k7_md_graph_n5retriever. Attention is 100 x (1 - grouped quality), so higher values identify larger observed improvement opportunities.

Answer outcomes

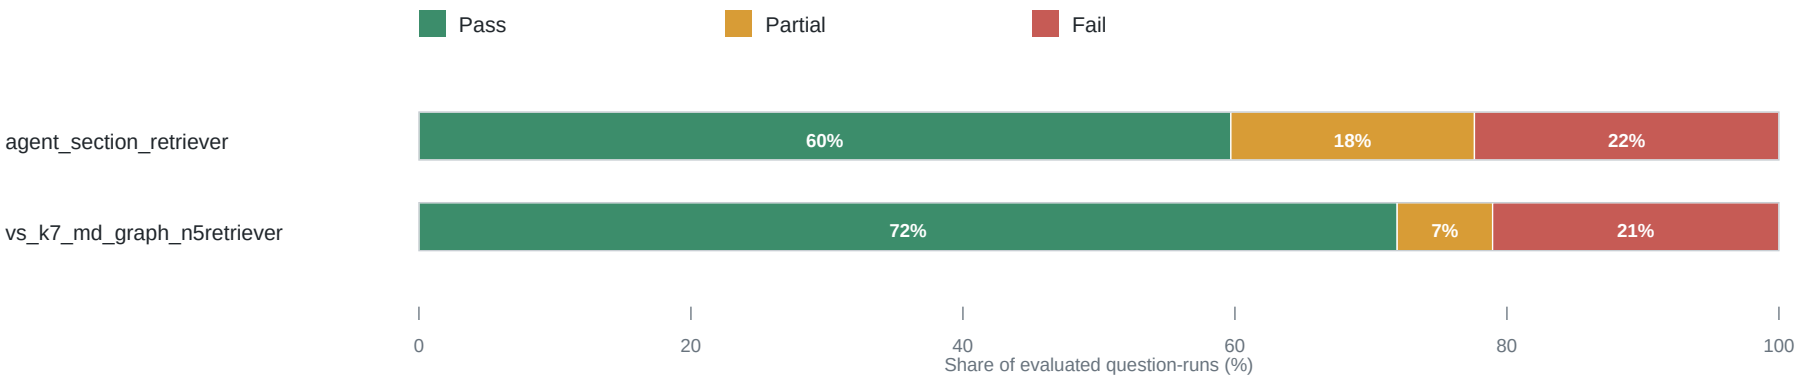

Configuration	Runs	Cases	Pass	Partial	Fail	Quality	Run SD	Article recall	Full source coverage	Mean context
agent_section_retriever	4	67	59.7%	17.9%	22.4%	0.687	0.057	90.7%	83.6%	15.4k
vs_k7_md_graph_n5retriever	3	57	71.9%	7.0%	21.1%	0.754	0.045	93.9%	89.5%	260.7k

Percentages aggregate all question-runs within a configuration. Run SD is the standard deviation of mean quality across repeated runs.

Retrieval diagnostics

Designated-source coverage

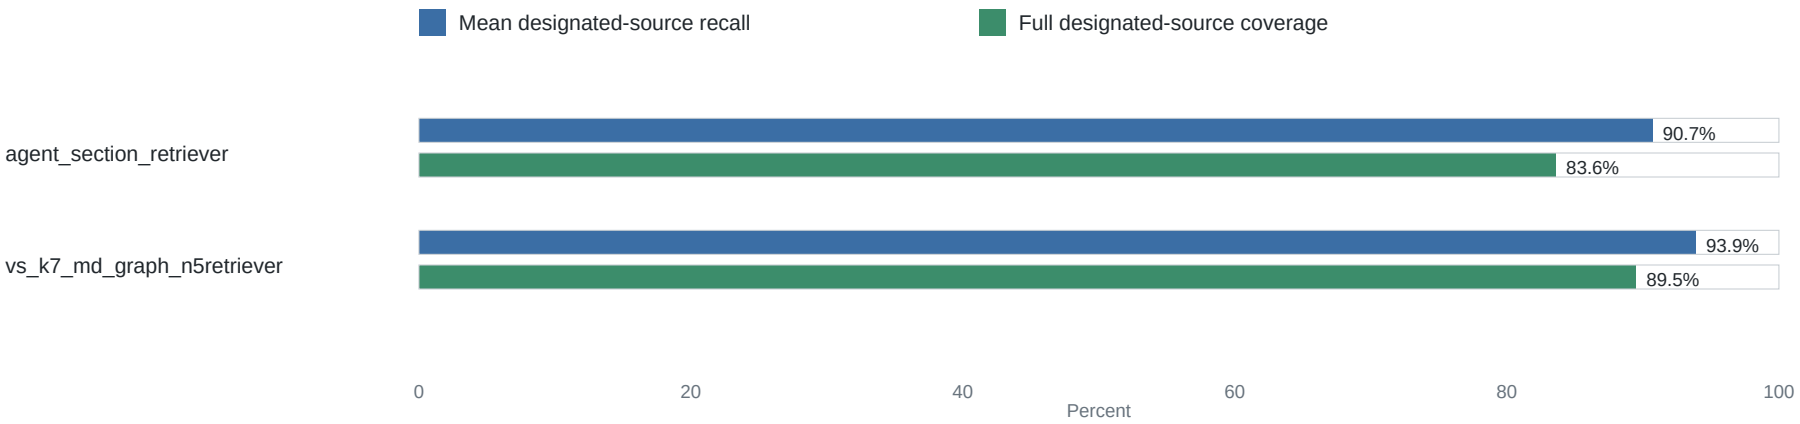

The analyzer recomputes recall after removing only final .txt and .html suffixes. This corrects historical .html.txt versus .html mismatches. Designated-source recall is diagnostic: alternative retrieved articles may still contain sufficient evidence.

Outcome and source-coverage combinations

Configuration	Pass, full	Pass, incomplete	Non-pass, full	Non-pass, missing
agent_section_retriever	33	7	23	4
vs_k7_md_graph_n5retriever	38	3	13	3

A non-pass result with full designated-source recall points toward answer synthesis, context noise, table interpretation, or judge behavior. A non-pass result with a missing designated source is consistent with a retrieval recall problem, but manual review is still required.

Per-question comparison

Cells show the mean quality score for each question and configuration across repeated runs.

	agent_section_retriever	vs_k7_md_graph_n5retriever
GT01	**0.75**	**1.00**
GT02	**1.00**	**1.00**
GT03	**0.00**	**0.00**
GT04	**1.00**	**1.00**
GT05	**1.00**	**1.00**
GT06	**0.50**	**1.00**
GT07	**0.50**	**0.33**
GT08	**0.00**	**0.67**
GT09	**0.88**	**1.00**
GT10	**0.00**	**0.00**
GT11	**0.12**	**0.00**
GT12	**0.83**	**1.00**
GT13	**1.00**	**1.00**
GT14	**1.00**	**1.00**
GT15	**1.00**	**1.00**
GT16	**0.75**	**0.50**
GT17	**1.00**	**1.00**
GT18	**1.00**	**1.00**
GT19	**0.75**	**0.83**

Fail (0) Partial (0.5) Pass (1)

Question	agent_section_retriever	vs_k7_md_graph_n5retriever
GT01	0.750	1.000
GT02	1.000	1.000
GT03	0.000	0.000
GT04	1.000	1.000
GT05	1.000	1.000
GT06	0.500	1.000
GT07	0.500	0.333
GT08	0.000	0.667
GT09	0.875	1.000
GT10	0.000	0.000
GT11	0.125	0.000
GT12	0.833	1.000
GT13	1.000	1.000
GT14	1.000	1.000
GT15	1.000	1.000
GT16	0.750	0.500
GT17	1.000	1.000
GT18	1.000	1.000
GT19	0.750	0.833

Per-question article recall

Fraction of each question's designated source articles that appeared in the retrieved set (retriever configurations only). 1.00 = every gold article retrieved; lower values mark a retrieval miss — usually a multi-source question.

Question	agent_section_retriever	vs_k7_md_graph_n5retriever
GT01	100.0%	100.0%

Question	agent_section_retriever	vs_k7_md_graph_n5retriever
GT02	100.0%	100.0%
GT03	58.3%	33.3%
GT04	100.0%	100.0%
GT05	100.0%	100.0%
GT06	100.0%	100.0%
GT07	100.0%	100.0%
GT08	100.0%	100.0%
GT09	25.0%	50.0%
GT10	100.0%	100.0%
GT11	100.0%	100.0%
GT12	88.9%	100.0%
GT13	100.0%	100.0%
GT14	68.8%	100.0%
GT15	100.0%	100.0%
GT16	100.0%	100.0%
GT17	100.0%	100.0%
GT18	100.0%	100.0%
GT19	100.0%	100.0%

Performance by task type

The 18 original category labels are consolidated into five task families so each group contains multiple questions. The attention score ranks observed improvement opportunities for the strongest retrieval configuration; it is not a statistical significance test.

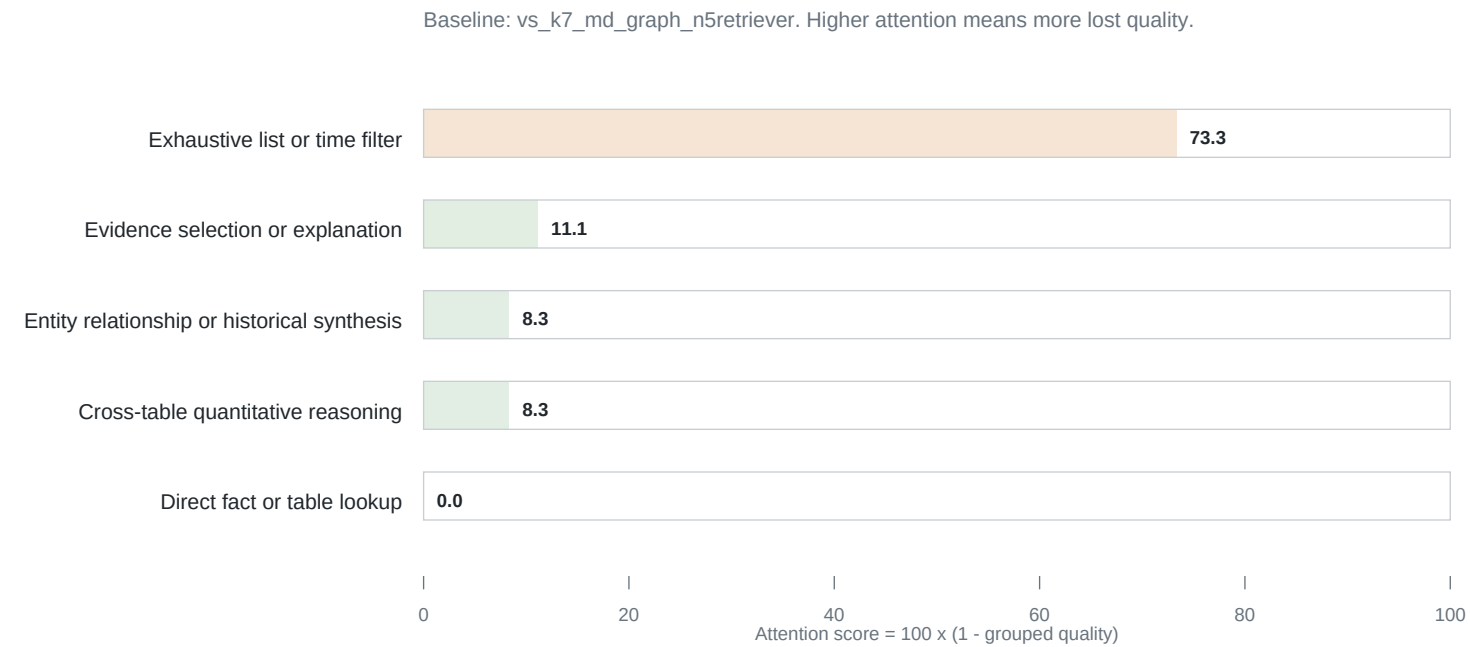

Task family	Questions	Baseline quality	Attention	Fail	Oracle quality	Oracle gap	Article recall
Exhaustive list or time filter	5	0.267	73.3	66.7%	n/a	n/a	86.7%
Evidence selection or explanation	3	0.889	11.1	11.1%	n/a	n/a	100.0%
Entity relationship or historical synthesis	6	0.917	8.3	5.6%	n/a	n/a	100.0%
Cross-table quantitative reasoning	2	0.917	8.3	0.0%	n/a	n/a	75.0%
Direct fact or table lookup	3	1.000	0.0	0.0%	n/a	n/a	100.0%

Baseline quality and failure rate use vs_k7_md_graph_n5retriever. Oracle gap equals oracle quality minus baseline quality: a large positive gap is consistent with a retrieval or context-selection opportunity, while low oracle quality indicates that the difficulty persists even when the designated articles are supplied.

Performance by source scope

Source scope groups questions by the number of designated articles. It tests whether the number of sources, rather than the reasoning operation, explains performance.

	agent_section_retriever	vs_k7_md_graph_n5retriever
1 article	**0.646**	**0.727**
2 articles	**0.867**	**0.867**
3+ articles	**0.591**	**0.667**

Source scope	Configuration	Questions	Question-runs	Pass	Partial	Fail	Quality	Article recall
1 article	agent_section_retriever	11	41	56.1%	17.1%	26.8%	0.646	100.0%
2 articles	agent_section_retriever	5	15	73.3%	26.7%	0.0%	0.867	80.0%
3+ articles	agent_section_retriever	3	11	54.5%	9.1%	36.4%	0.591	70.5%
1 article	vs_k7_md_graph_n5retriever	11	33	69.7%	6.1%	24.2%	0.727	100.0%
2 articles	vs_k7_md_graph_n5retriever	5	15	80.0%	13.3%	6.7%	0.867	90.0%
3+ articles	vs_k7_md_graph_n5retriever	3	9	66.7%	0.0%	33.3%	0.667	77.8%

Question-runs are repeated measurements of the listed unique questions. Source-scope groups remain small and mix different task families, so this view should be used as a diagnostic rather than a general performance estimate.

Run-to-run variability

Each dot is one run; the larger dot is the configuration mean.

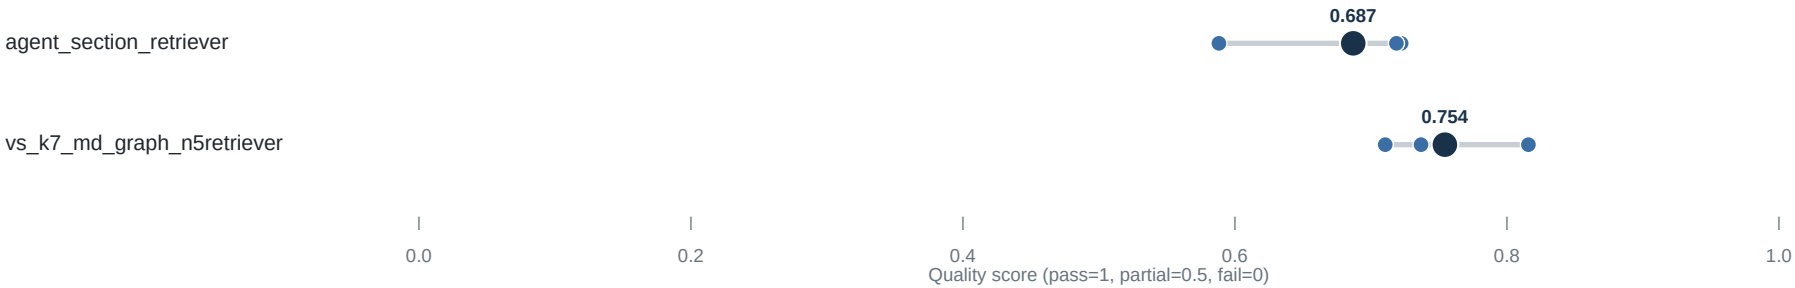

Run	Configuration	Cases	Pass	Partial	Fail	Quality	Article recall	Full source coverage
sleepy_diffie	agent_section_retriever	18	11	4	3	0.722	94.9%	88.9%
tender_brin	agent_section_retriever	16	11	1	4	0.719	85.9%	75.0%
jolly_minsky	agent_section_retriever	16	10	3	3	0.719	90.1%	81.2%
wizardly_hopper	agent_section_retriever	17	8	4	5	0.588	91.2%	88.2%
bold_wozniak	vs_k7_md_graph_n5retriever	19	13	2	4	0.737	93.9%	89.5%
awesome_valiant	vs_k7_md_graph_n5retriever	19	13	1	5	0.711	93.9%	89.5%
reverent_wilkes	vs_k7_md_graph_n5retriever	19	15	1	3	0.816	93.9%	89.5%

Methodology and limitations

Scoring

Pass is assigned 1, partial is assigned 0.5, and fail is assigned 0. The quality score is the arithmetic mean across the applicable question-runs. Pass, partial, and fail percentages are also reported separately so the weighted score does not hide the underlying verdicts.

Task-family attention

The attention score is 100 multiplied by one minus grouped quality. A score of 66 means that the group lost 66% of the available quality points under the selected retrieval baseline. It is a prioritization aid, not an estimate of population-wide failure risk. Task-family and source-scope results must be read with their unique-question counts because the evaluation set contains only 19 questions.

Configuration grouping

Runs are grouped by a CSV experiment_name or config_name when one is present. If no explicit name is present, grouping falls back to evaluation mode, context source, retriever kind and label, Top K, hybrid weights, answer model, and judge model. RAGAS-generated filenames identify individual runs but do not define configurations.

Framework compatibility

Legacy oracle runs stored under pass_fail_only directories are excluded by default because they could not produce partial verdicts. Use --include-binary-oracle only for a labeled historical appendix, not for a direct three-level score comparison.

Interpretation limits

An LLM produced the answers and another LLM applied the rubric, so repeated runs can vary. Article recall measures coverage of designated ground-truth sources, not every possible supporting source. Full article recall does not prove that the model used the evidence correctly, and missing a designated source does not prove the context was insufficient when another article contains equivalent evidence. Differences between configurations should be interpreted together with run counts and per-question behavior.

Included experiment files

Run file
experiments/awesome_valiant.csv
experiments/bold_wozniak.csv
experiments/jolly_minsky.csv
experiments/reverent_wilkes.csv
experiments/sleepy_diffie.csv
experiments/tender_brin.csv
experiments/wizardly_hopper.csv

Excluded files

Run file	Reason
datasets/wikipedia_GT_agentic_section_inputs.csv	missing columns: ['score']