GT baseline: 6.1 survivor against 7.1 final

Generated 2026-10-04 08:31 EDT. Sources: wikipedia_retrieval_engine/Capstone Checkpoint 7.1/cost_efficiency/wikipedia_GT

Question set	Configurations	Runs	Question-runs	Baseline	Highest quality	Lowest list cost	Fastest
GT test set	2	6	114	6.1 survivor	6.1 survivor (0.789)	6.1 survivor ($0.0485)	6.1 survivor (13.1 s)

How to read this report

Quality. pass = 1, partial = 0.5, fail = 0, averaged over question-runs (RAGAS judge).

List cost. every input token at the full price: the comparison of record. Billed cost charges cached input at the cached rate, so it falls when runs repeat the same questions.

Seconds. agent graph plus answer per question, judge excluded. Provider speed varies by hour; §6 shows how much.

Dominated. another configuration is at least as good and no more expensive (or slower), and better on one of the two.

Paired. each question's mean over its runs, configuration minus baseline; p from the Wilcoxon signed-rank test over the questions that changed.

Paths. Table: the TOC review flagged a table. Text, escalated: an assessment asked for more evidence. Text, ready: neither.

GT test set

2 configuration(s), 6 run(s), 114 question-runs. Baseline: 6.1 survivor (requested).

1. Frontier: quality against cost and time

Configuration	Runs (with telemetry)	Quality (run range)	Pass / partial / fail	List $/q (run range)	Billed $/q	Cached	Seconds mean / median	Calls/q	Input tokens/q	List $ per quality point	Dominated on cost by	Dominated on time by
6.1 survivor	3 (3)	0.789 (0.737-0.816)	40/10/7	$0.0485 ($0.0402-$0.0570)	$0.0350	35%	13.1 / 11.8	6.77	56.8K	$0.0615	-	-
7.1 final	3 (3)	0.711 (0.658-0.763)	37/7/13	$0.0502 ($0.0459-$0.0539)	$0.0348	38%	13.8 / 12.6	6.84	59.3K	$0.0706	6.1 survivor	6.1 survivor

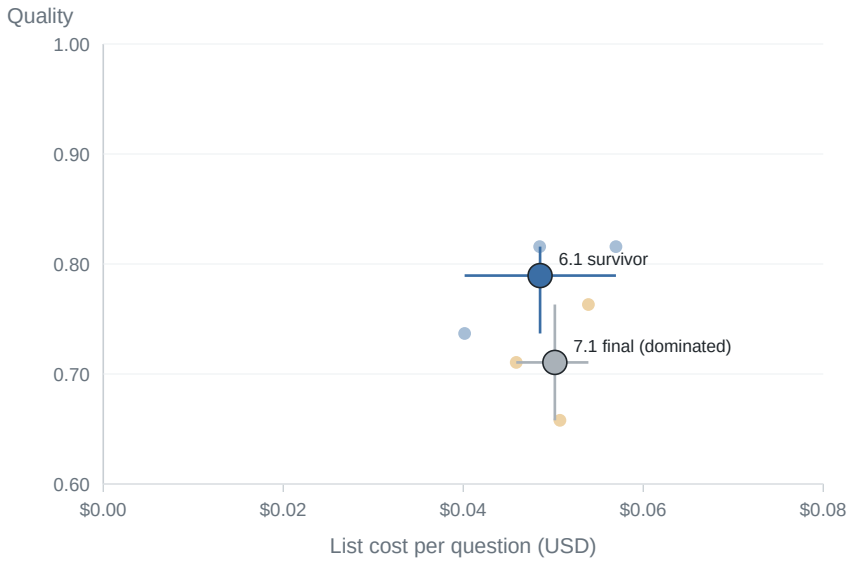

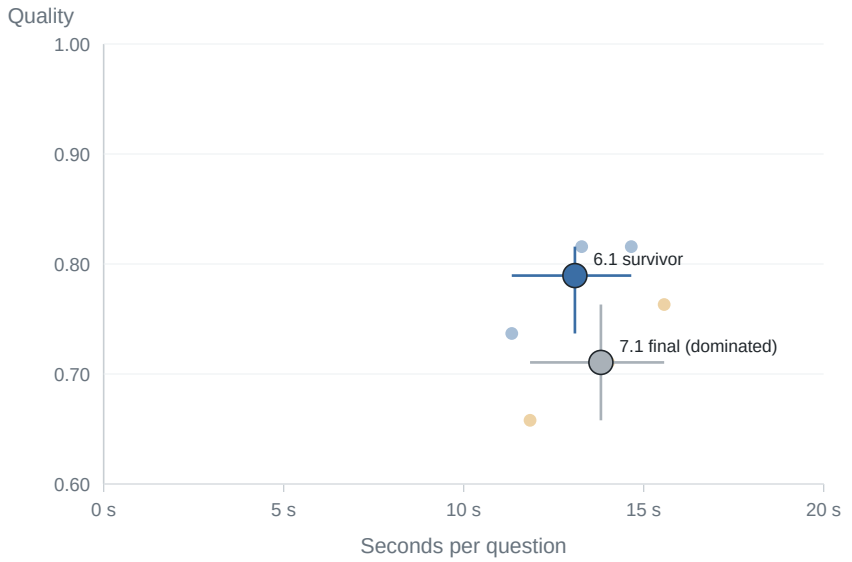

Small dots are runs; a large dot is the configuration mean, with lines across its run range. Gray: dominated.

2. Paired against the baseline (6.1 survivor)

Configuration	Questions	Quality	Δ quality	Better / worse / same	p (quality)	Δ list $/q	p (cost)	Δ seconds	p (time)	Trade-off
7.1 final	19	0.789 -> 0.711	-0.079	2 / 6 / 11	0.188	+$0.0016 (+3%)	0.679	+0.7	0.061	dominated by the baseline: no better, no cheaper

Per-question change in quality, worst first (right column: change in list cost):

7.1 final against 6.1 survivor

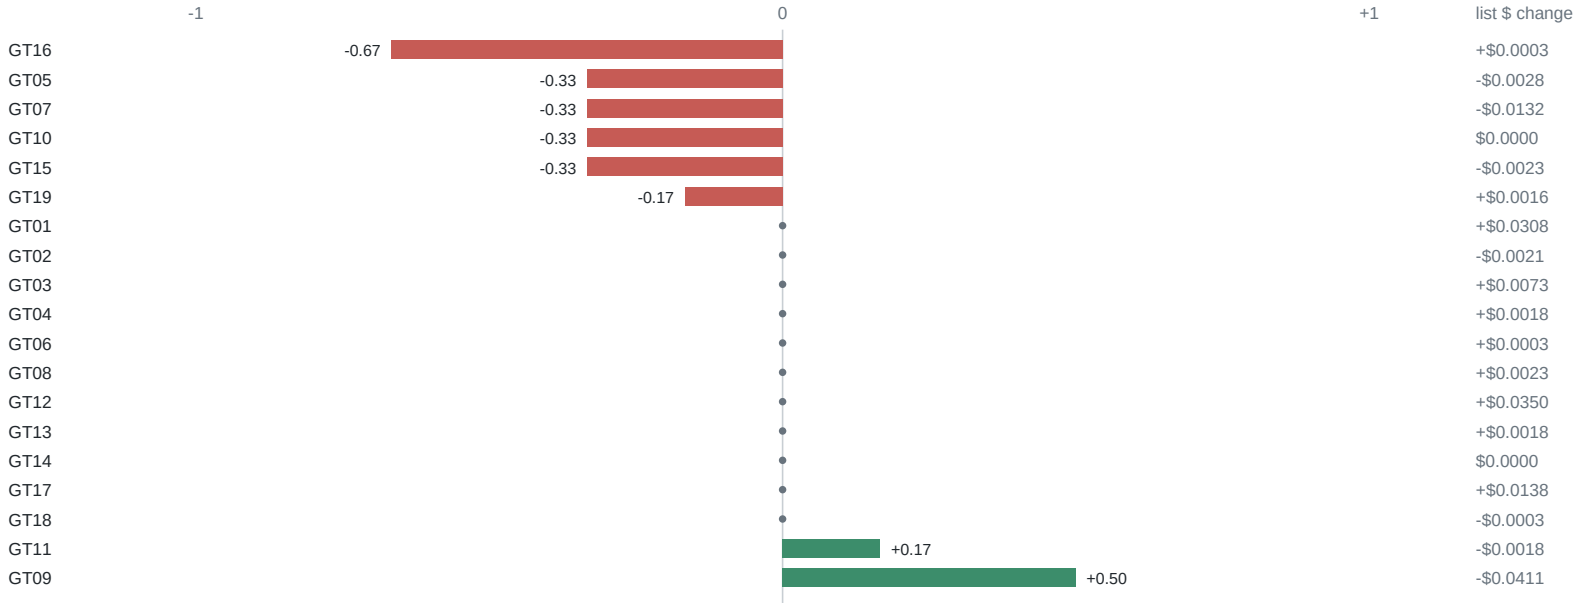

3. Where the cost goes: by role

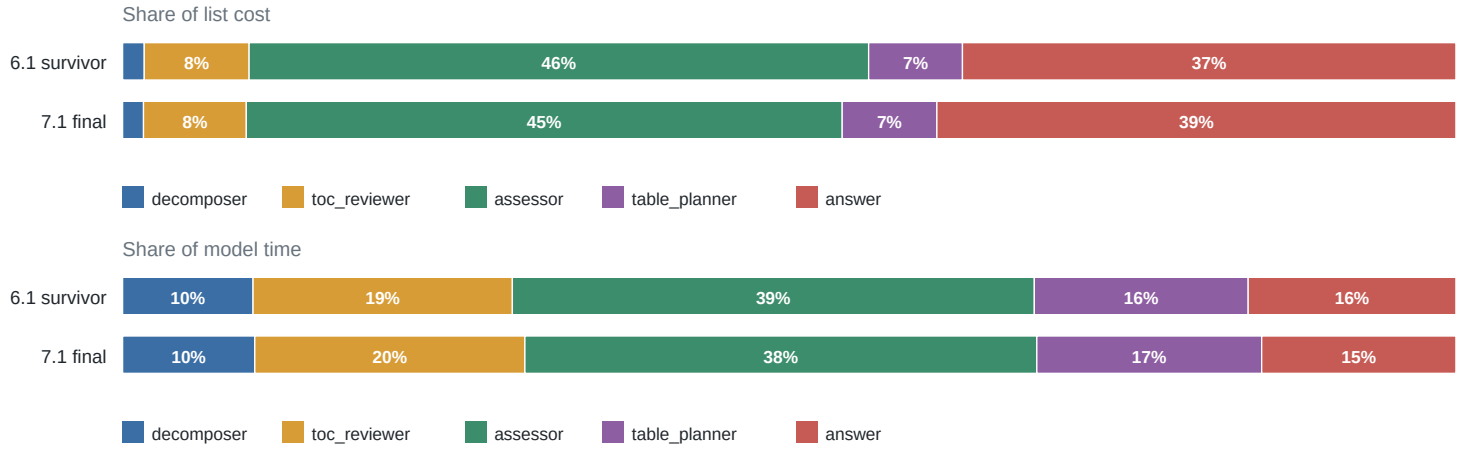

Configuration	Role	Calls/q	Input tokens/q	Output tokens/q	Cached	List $/q	Share of list	Model s/q	Share of time	Median s/call
6.1 survivor	decomposer	1.00	0.4K	105	0%	$0.0008	2%	1.0	10%	1.02
6.1 survivor	toc_reviewer	1.00	3.0K	353	55%	$0.0038	8%	2.1	19%	1.99
6.1 survivor	assessor	2.09	26.9K	534	38%	$0.0226	46%	4.2	39%	1.85
6.1 survivor	table_planner	1.68	3.7K	137	38%	$0.0034	7%	1.7	16%	0.95
6.1 survivor	answer	1.00	22.8K	184	30%	$0.0180	37%	1.7	16%	1.42
7.1 final	decomposer	1.00	0.4K	104	0%	$0.0008	2%	1.1	10%	1.04
7.1 final	toc_reviewer	1.00	3.0K	361	82%	$0.0039	8%	2.3	20%	2.20
7.1 final	assessor	2.02	26.9K	505	34%	$0.0224	45%	4.3	38%	1.89
7.1 final	table_planner	1.82	3.9K	140	47%	$0.0036	7%	1.9	17%	1.00
7.1 final	answer	1.00	25.1K	149	37%	$0.0195	39%	1.6	15%	1.52

4. By path: text ready, text escalated, table

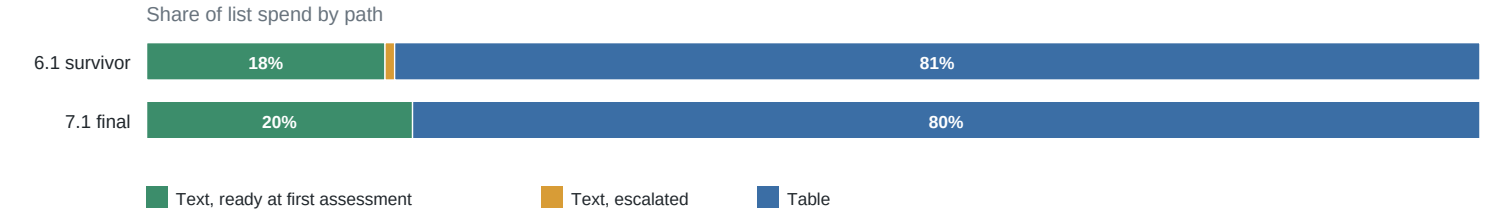

Configuration	Path	Question-runs (share)	Quality	List $/q	Billed $/q	Seconds/q	Share of list spend
6.1 survivor	Text, ready at first assessment	13 (23%)	0.962	$0.0381	$0.0203	8.9	18%
6.1 survivor	Text, escalated	1 (2%)	1.000	$0.0200	$0.0200	8.2	1%
6.1 survivor	Table	43 (75%)	0.733	$0.0524	$0.0398	14.5	81%
7.1 final	Text, ready at first assessment	15 (26%)	0.833	$0.0380	$0.0191	9.3	20%
7.1 final	Table	42 (74%)	0.667	$0.0545	$0.0404	15.4	80%

5. Tail and waste

Configuration	List p50	List p90	List max	Max / mean	Seconds p50	Seconds p90	Seconds max	Spend on fails	Spend on partials	List $ per pass
6.1 survivor	$0.0333	$0.0813	$0.3144	6.5x	11.8	20.1	30.5	13%	34%	$0.0692
7.1 final	$0.0343	$0.0902	$0.2360	4.7x	12.6	21.2	30.8	23%	14%	$0.0773

Most expensive question-runs:

Configuration	Run	Question	Path	Verdict	List $	Seconds	Calls
6.1 survivor	hardcore_page	GT09	Table	partial	$0.3144	20.4	9
7.1 final	confident_zuckerberg	GT09	Table	pass	$0.2360	23.3	8
7.1 final	admiring_diffie	GT01	Table	pass	$0.1740	17.0	11
6.1 survivor	confident_stallman	GT09	Table	partial	$0.1444	17.0	7
7.1 final	eloquent_hoare	GT12	Table	pass	$0.1207	25.2	11
6.1 survivor	confident_stallman	GT07	Table	partial	$0.1064	21.3	10
6.1 survivor	confident_stallman	GT03	Table	fail	$0.0938	26.4	8
7.1 final	eloquent_hoare	GT03	Table	fail	$0.0932	21.4	8
7.1 final	admiring_diffie	GT09	Table	pass	$0.0931	17.6	9
7.1 final	confident_zuckerberg	GT12	Table	pass	$0.0903	24.8	10

6. Comparability: run order, caching, and provider speed

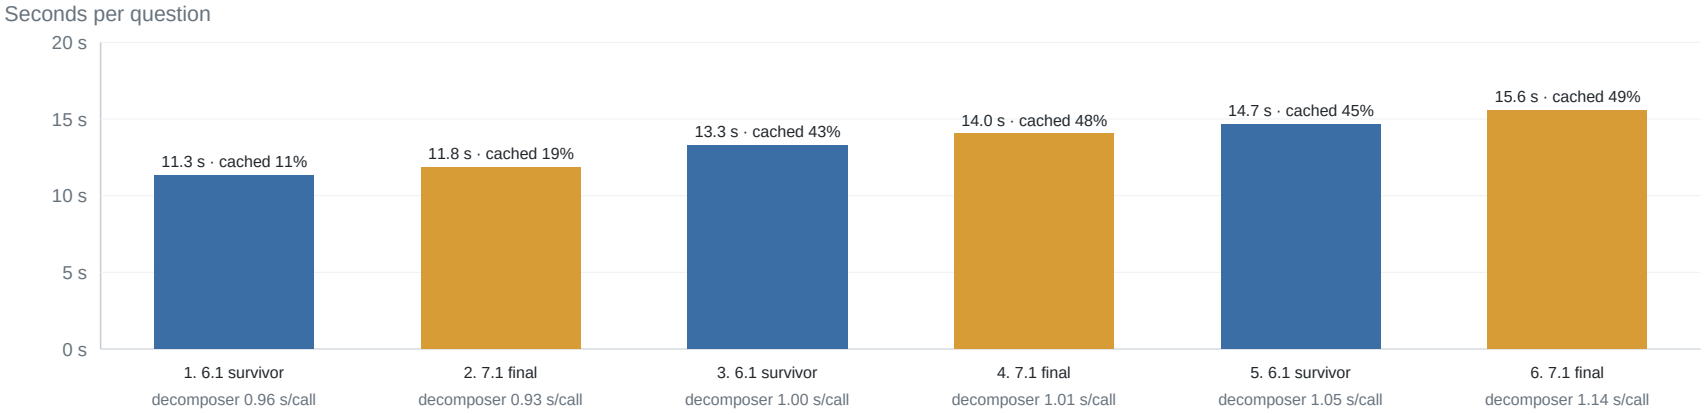

Run (order)	Results CSV	Finished	Quality	List $/q	Billed $/q	Billed / list	Cached	Seconds per q	Decomp. s/call	TOC s/call	Assessor s/call	Planner s/call	Answer s/call
1. 6.1 survivor	eager_torvalds	2026-10-04 01:41	0.737	$0.0402	$0.0368	0.92	11%	11.3	0.96	1.80	1.58	0.93	1.30
2. 7.1 final	admiring_diffie	2026-10-04 01:45	0.658	$0.0507	$0.0432	0.85	19%	11.8	0.93	1.80	1.64	0.92	1.46
3. 6.1 survivor	hardcore_page	2026-10-04 01:50	0.816	$0.0570	$0.0372	0.65	43%	13.3	1.00	1.76	1.82	0.94	1.32
4. 7.1 final	eloquent_hoare	2026-10-04 01:55	0.711	$0.0459	$0.0286	0.62	48%	14.0	1.01	2.13	1.91	0.98	1.52
5. 6.1 survivor	confident_stallman	2026-10-04 02:00	0.816	$0.0485	$0.0310	0.64	45%	14.7	1.05	2.41	2.13	1.09	1.64
6. 7.1 final	confident_zuckerberg	2026-10-04 02:06	0.763	$0.0539	$0.0325	0.60	49%	15.6	1.14	2.52	2.08	1.10	1.63

Run times from: run_manifest.csv.

Warnings

- 6.1 survivor: list cost per question ranges from $0.0402 to $0.0570 across its own runs (42% apart).

- Cached share ranges from 11% to 49% across runs (it grows with repeats of the same questions). Billed cost depends on run order; compare list cost.

7. Safeguard markers

Configuration	Question-runs	Evidence notice	Noticed: pass / partial / fail	Grounding checked	Flagged	Flagged: pass / partial / fail	Check passed, judge failed	Input sanitized
6.1 survivor	57	17 (30%)	12/4/1	not checked	-	-	-	0
7.1 final	57	17 (30%)	11/4/2	57	10 (18%)	8/1/1	12	0

A flagged answer the judge passed is a false alarm the user still sees (policy 'surface' labels it). 'Check passed, judge failed' counts wrong answers the deterministic check cannot see.

Spans the grounding check could not find in the evidence:

Span	Kind	Times	Judge: pass / partial / fail	Configurations
Russia-related Olympic	name	3	3/0/0	7.1 final
Russian Empire	name	2	2/0/0	7.1 final
Winter Olympic Games Sport	name	2	2/0/0	7.1 final
Spanish armies, led by Fernando Álvarez de Toledo, 3rd Duke of Alba, defending the claim of King Philip II of Spain to the Portuguese throne, defeat the armies of Portuguese claimant António, Prior of Crato	quotation	1	1/0/0	7.1 final
Summer Olympic Games Sport	name	1	0/0/1	7.1 final
To Barack Obama	name	1	0/1/0	7.1 final
To Bill Clinton	name	1	0/1/0	7.1 final

8. Routing

Not requested (--route-cheap and --route-strong).

Route mix: no run on this set records a router decision (column 'route').