Survivor quality, cost, and latency

Generated 2026-10-04 16:39 EDT. Sources: wikipedia_retrieval_engine/Capstone Checkpoint 7.1/final_validation/runs

Question set	Configurations	Runs	Question-runs	Baseline	Highest quality	Lowest list cost	Fastest
Validation set	3	9	225	6.1 survivor	7.1Final Survivor (0.933)	6.1 survivor ($0.0422)	6.1 survivor (12.8 s)

How to read this report

Quality. pass = 1, partial = 0.5, fail = 0, averaged over question-runs (RAGAS judge).

List cost. every input token at the full price: the comparison of record. Billed cost charges cached input at the cached rate, so it falls when runs repeat the same questions.

Seconds. agent graph plus answer per question, judge excluded. Provider speed varies by hour; §6 shows how much.

Dominated. another configuration is at least as good and no more expensive (or slower), and better on one of the two.

Paired. each question's mean over its runs, configuration minus baseline; p from the Wilcoxon signed-rank test over the questions that changed.

Paths. Table: the TOC review flagged a table. Text, escalated: an assessment asked for more evidence. Text, ready: neither.

Loader warnings

- Task-family analysis left these categories unclassified: Cross-article comparison, Diagnostic reasoning and limitations, Inference limitation, People connection: shared work and roles

Validation set

3 configuration(s), 9 run(s), 225 question-runs. Baseline: 6.1 survivor (requested).

1. Frontier: quality against cost and time

Configuration	Runs (with telemetry)	Quality (run range)	Pass / partial / fail	List $/q (run range)	Billed $/q	Cached	Seconds mean / median	Calls/q	Input tokens/q	List $ per quality point	Dominated on cost by	Dominated on time by
5.1 survivor	3 (1)	0.880 (0.860-0.900)	62/8/5	$0.0431 ($0.0431)	$0.0356	23%	16.1 / 14.1	6.72	49.0K	$0.0501	6.1 survivor	6.1 survivor
6.1 survivor	3 (3)	0.860 (0.800-0.900)	60/9/6	$0.0422 ($0.0358-$0.0521)	$0.0296	38%	12.8 / 11.0	5.79	49.1K	$0.0491	-	-
7.1Final Survivor	3 (3)	0.933 (0.920-0.940)	67/6/2	$0.1068 ($0.0962-$0.1253)	$0.0751	44%	18.7 / 14.0	6.55	53.8K	$0.1144	-	-

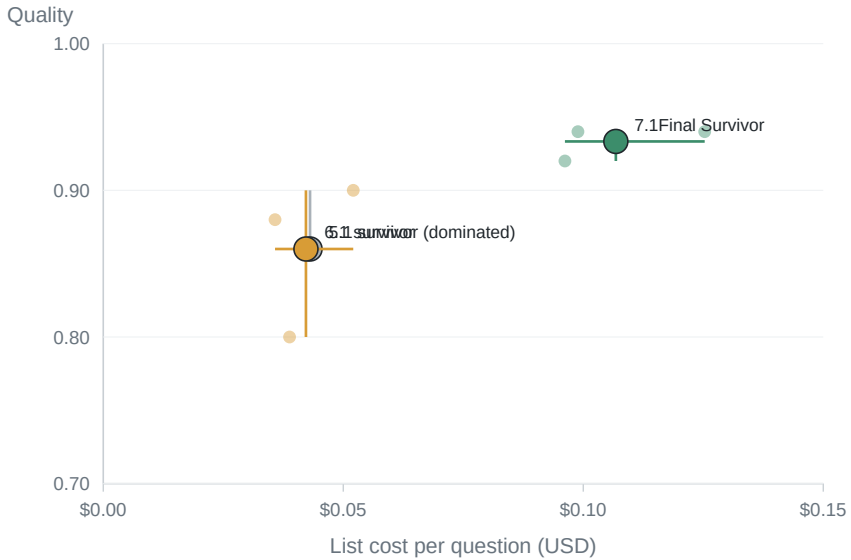

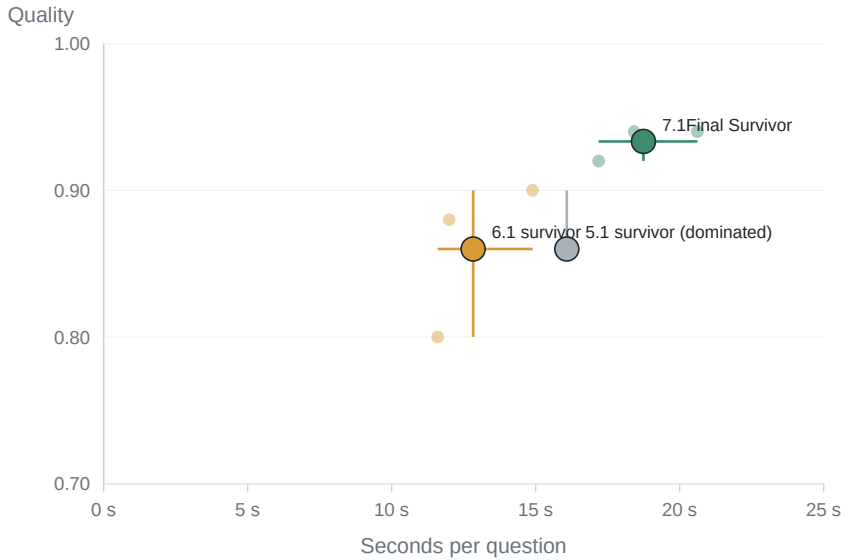

Small dots are runs; a large dot is the configuration mean, with lines across its run range. Gray: dominated.

2. Paired against the baseline (6.1 survivor)

Configuration	Questions	Quality	Δ quality	Better / worse / same	p (quality)	Δ list $/q	p (cost)	Δ seconds	p (time)	Trade-off
5.1 survivor	25	0.860 -> 0.880	+0.020	4 / 3 / 18	0.688	+$0.0009 (+2%)	0.895	+3.3	<0.001	pays $0.0043 per question for each +0.10 quality
7.1Final Survivor	25	0.860 -> 0.933	+0.073	6 / 1 / 18	0.094	+$0.0646 (+153%)	<0.001	+5.9	0.002	pays $0.0880 per question for each +0.10 quality

Per-question change in quality, worst first (right column: change in list cost):

5.1 survivor against 6.1 survivor

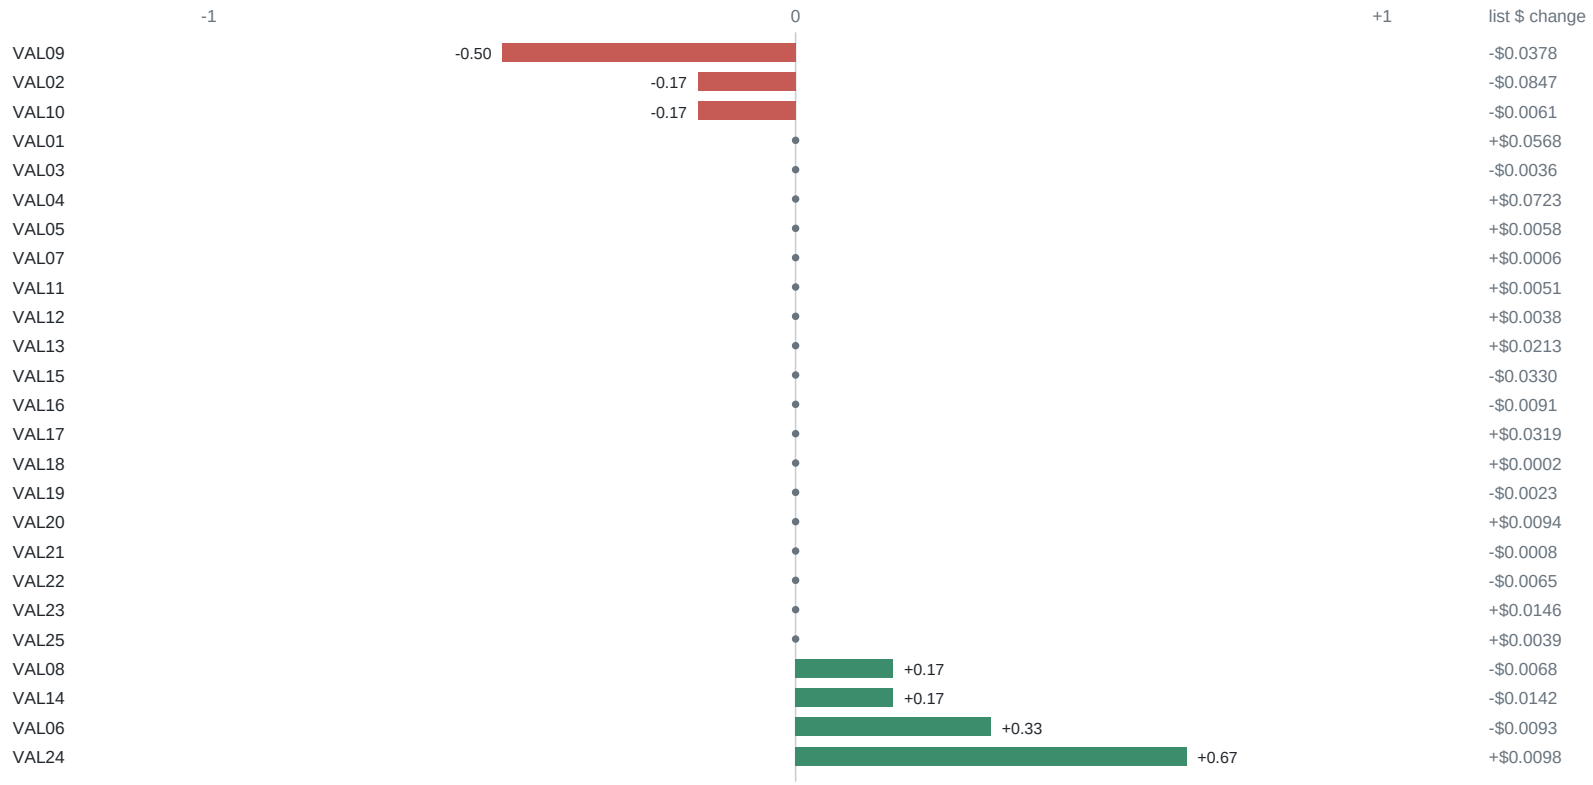

7.1Final Survivor against 6.1 survivor

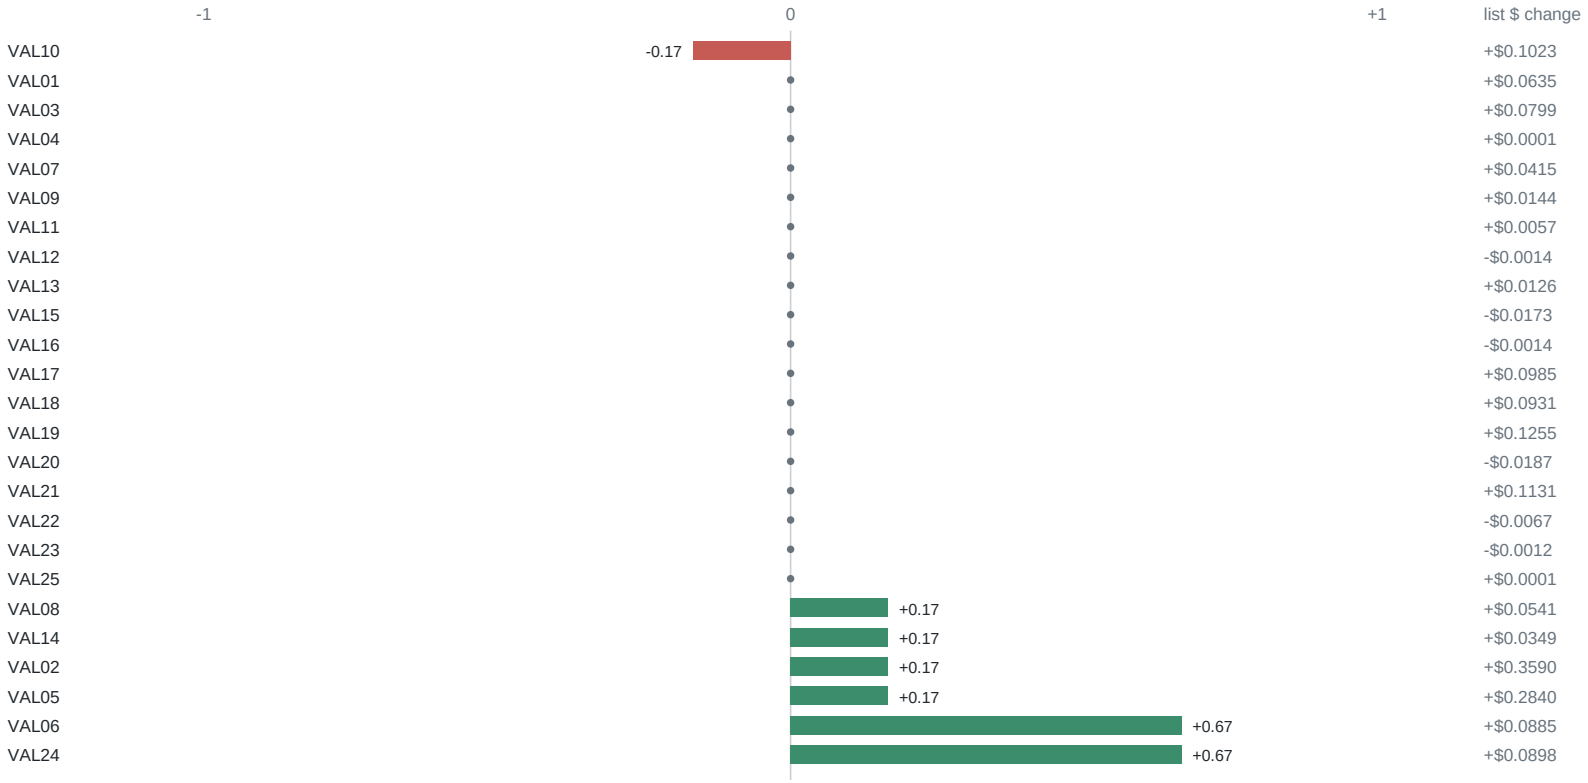
-1
0
+1
list $ change
VAL10
-0.17
+$0.1023
VAL01
+$0.0635
VAL03
+$0.0799
VAL04
+$0.0001
VAL07
+$0.0415
VAL09
+$0.0144
VAL11
+$0.0057
VAL12
-$0.0014
VAL13
+$0.0126
VAL15
-$0.0173
VAL16
-$0.0014
VAL17
+$0.0985
VAL18
+$0.0931
VAL19
+$0.1255
VAL20
-$0.0187
VAL21
+$0.1131
VAL22
-$0.0067
VAL23
-$0.0012
VAL25
+$0.0001
VAL08
+0.17
+$0.0541
VAL14
+0.17
+$0.0349
VAL02
+0.17
+$0.3590
VAL05
+0.17
+$0.2840
VAL06
+0.67
+$0.0885
VAL24
+0.67
+$0.0898

3. Where the cost goes: by role

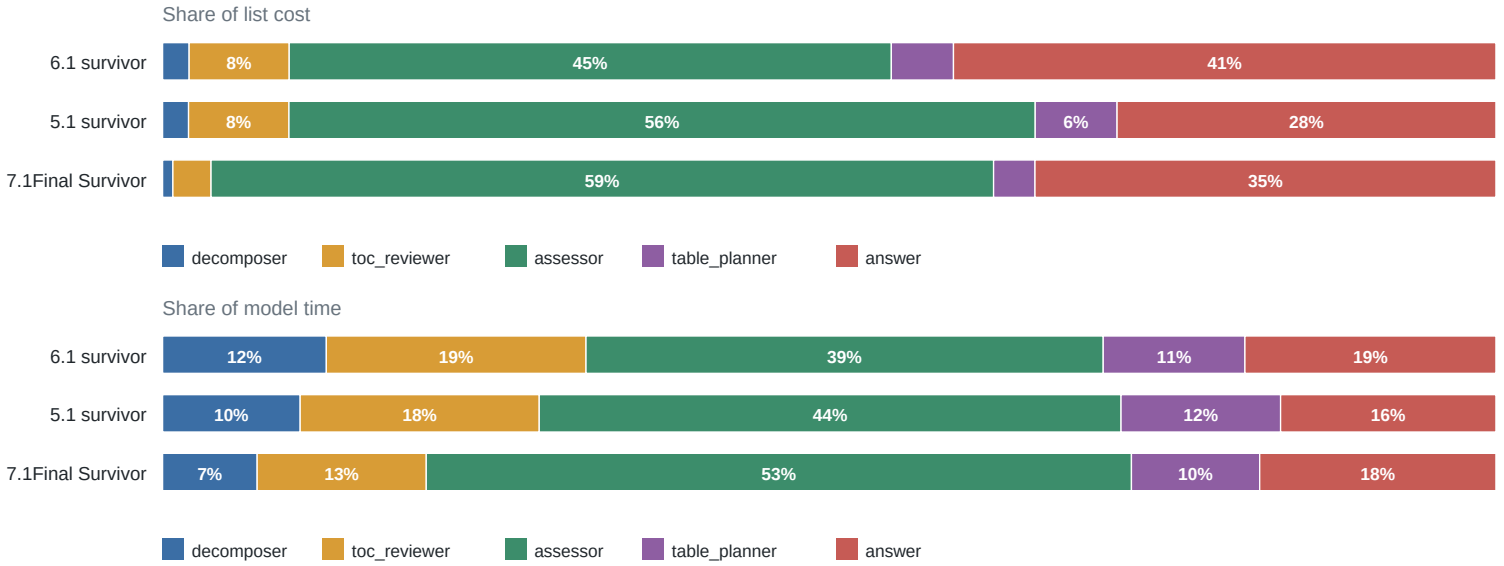

Configuration	Role	Calls/q	Input tokens/q	Output tokens/q	Cached	List $/q	Share of list	Model s/q	Share of time	Median s/call
6.1 survivor	decomposer	1.00	0.4K	115	0%	$0.0008	2%	1.2	12%	1.22
6.1 survivor	toc_reviewer	1.00	2.4K	310	48%	$0.0032	8%	2.0	19%	1.92
6.1 survivor	assessor	1.85	22.4K	504	36%	$0.0191	45%	3.9	39%	1.94
6.1 survivor	table_planner	0.93	2.2K	73	32%	$0.0020	5%	1.1	11%	1.08
6.1 survivor	answer	1.00	21.7K	205	40%	$0.0172	41%	1.9	19%	1.82
5.1 survivor	decomposer	1.00	0.4K	115	0%	$0.0008	2%	1.3	10%	1.32
5.1 survivor	toc_reviewer	1.00	2.4K	321	0%	$0.0032	8%	2.3	18%	2.25
5.1 survivor	assessor	2.48	28.1K	681	39%	$0.0241	56%	5.6	44%	2.06
5.1 survivor	table_planner	1.24	2.9K	106	5%	$0.0026	6%	1.5	12%	1.10
5.1 survivor	answer	1.00	15.2K	189	1%	$0.0122	28%	2.1	16%	1.96
7.1Final Survivor	decomposer	1.00	0.4K	112	0%	$0.0008	1%	1.1	7%	1.11
7.1Final Survivor	toc_reviewer	1.00	2.3K	296	78%	$0.0031	3%	2.0	13%	2.02
7.1Final Survivor	assessor	2.09	25.1K	689	36%	$0.0627	59%	8.5	53%	4.06
7.1Final Survivor	table_planner	1.45	3.7K	120	38%	$0.0033	3%	1.5	10%	0.99
7.1Final Survivor	answer	1.00	22.3K	218	51%	$0.0369	35%	2.8	18%	2.13

4. By path: text ready, text escalated, table

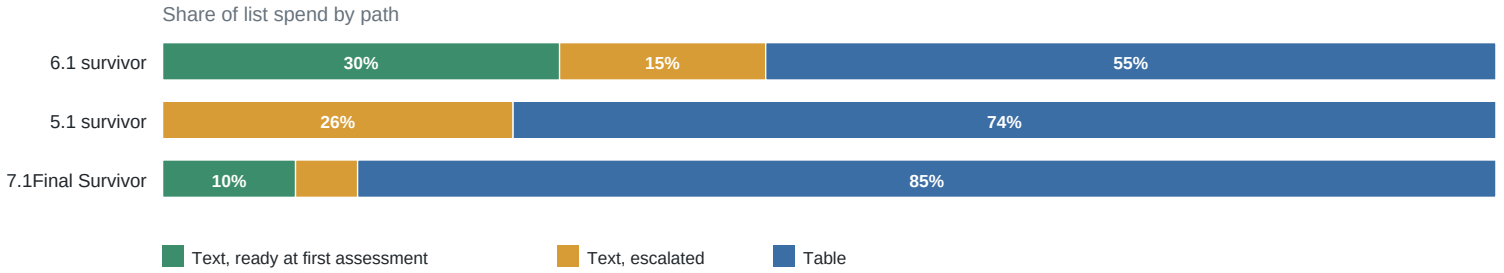

Configuration	Path	Question-runs (share)	Quality	List $/q	Billed $/q	Seconds/q	Share of list spend
6.1 survivor	Text, ready at first assessment	34 (45%)	1.000	$0.0277	$0.0184	9.5	30%
6.1 survivor	Text, escalated	5 (7%)	0.800	$0.0979	$0.0579	17.5	15%
6.1 survivor	Table	36 (48%)	0.736	$0.0482	$0.0363	15.3	55%
5.1 survivor	Text, escalated	35 (47%)	0.929	$0.0283	$0.0229	14.9	26%
5.1 survivor	Table	40 (53%)	0.838	$0.0529	$0.0440	16.9	74%
7.1Final Survivor	Text, ready at first assessment	30 (40%)	1.000	$0.0266	$0.0125	9.3	10%
7.1Final Survivor	Text, escalated	6 (8%)	1.000	$0.0621	$0.0393	13.6	5%
7.1Final Survivor	Table	39 (52%)	0.872	$0.1753	$0.1288	26.8	85%

5. Tail and waste

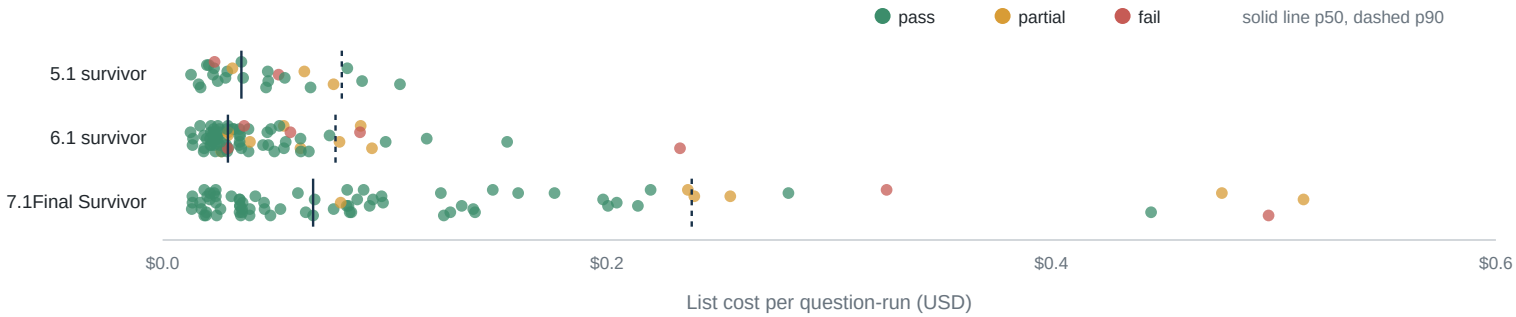

Configuration	List p50	List p90	List max	Max / mean	Seconds p50	Seconds p90	Seconds max	Spend on fails	Spend on partials	List $ per pass
6.1 survivor	$0.0294	$0.0778	$0.2328	5.5x	11.0	19.3	35.6	15%	16%	$0.0528
5.1 survivor	$0.0355	$0.0806	$0.1068	2.5x	14.1	22.2	30.1	7%	16%	$0.0538
7.1Final Survivor	$0.0677	$0.2381	$0.5134	4.8x	14.0	35.7	46.8	10%	22%	$0.1195

Most expensive question-runs:

Configuration	Run	Question	Path	Verdict	List $	Seconds	Calls
7.1Final Survivor	xenodochial_tarjan	VAL02	Table	partial	$0.5134	39.3	12
7.1Final Survivor	sweet_dijkstra	VAL02	Table	fail	$0.4977	37.6	17
7.1Final Survivor	thirsty_stallman	VAL02	Table	partial	$0.4767	31.7	11
7.1Final Survivor	sweet_dijkstra	VAL05	Table	pass	$0.4448	41.1	15
7.1Final Survivor	thirsty_stallman	VAL05	Table	fail	$0.3258	35.8	14
7.1Final Survivor	sweet_dijkstra	VAL19	Table	pass	$0.2816	28.5	8
7.1Final Survivor	xenodochial_tarjan	VAL05	Table	partial	$0.2555	40.3	17
7.1Final Survivor	thirsty_stallman	VAL10	Table	partial	$0.2393	46.8	10
7.1Final Survivor	sweet_dijkstra	VAL10	Table	partial	$0.2364	46.0	10
6.1 survivor	angry_babbage	VAL02	Text, escalated	fail	$0.2328	35.6	14

6. Comparability: run order, caching, and provider speed

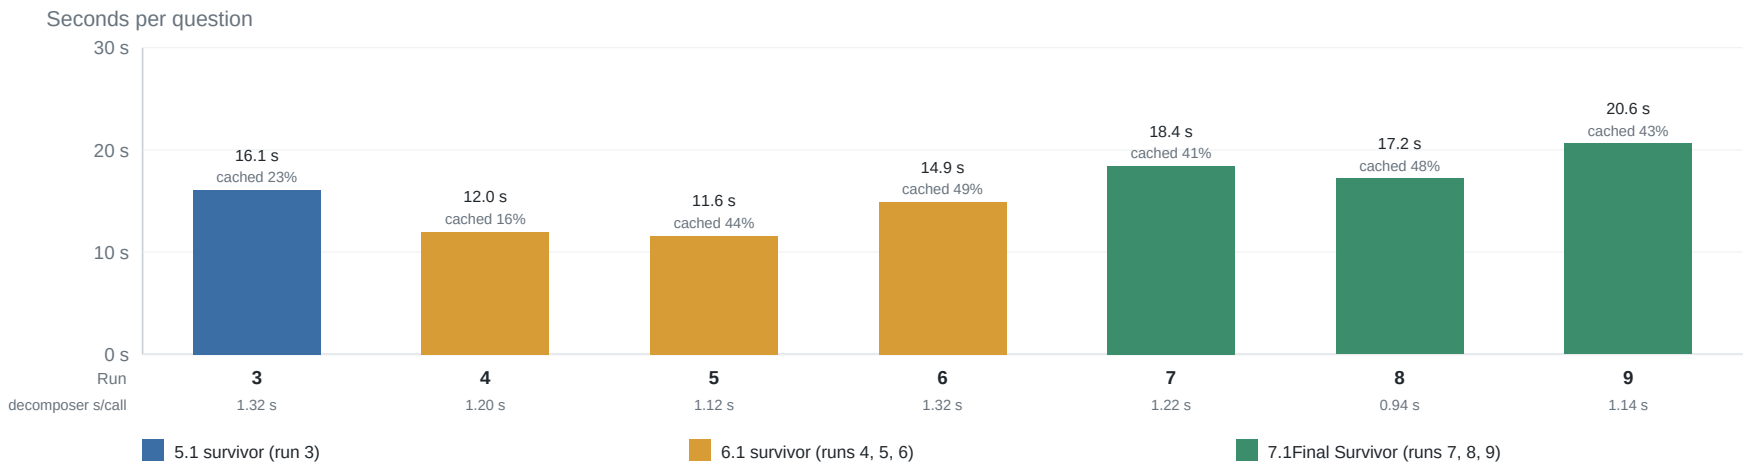

Run (order)	Results CSV	Finished	Quality	List $/q	Billed $/q	Billed / list	Cached	Seconds per q	Decomp. s/call	TOC s/call	Assessor s/call	Planner s/call	Answer s/call
1. 5.1 survivor	eager_rossum	2026-09-23 12:05	0.900	n/a	n/a	n/a	n/a	n/a	n/a	n/a	n/a	n/a	n/a
2. 5.1 survivor	objective_hopper	2026-09-23 12:17	0.880	n/a	n/a	n/a	n/a	n/a	n/a	n/a	n/a	n/a	n/a
3. 5.1 survivor	quirky_adleman	2026-09-27 13:30	0.860	$0.0431	$0.0356	0.83	23%	16.1	1.32	2.25	2.06	1.10	1.96
4. 6.1 survivor	dazzling_minsky	2026-09-27 15:48	0.880	$0.0358	$0.0314	0.88	16%	12.0	1.20	1.84	1.89	1.08	1.82
5. 6.1 survivor	nifty_mccarthy	2026-09-27 15:54	0.800	$0.0388	$0.0256	0.66	44%	11.6	1.12	1.67	1.66	1.07	1.65
6. 6.1 survivor	angry_babbage	2026-09-27 16:00	0.900	$0.0521	$0.0318	0.61	49%	14.9	1.32	2.22	2.23	1.08	2.04
7. 7.1Final Survivor	xenodochial_tarjan	2026-10-04 15:56	0.940	$0.0989	$0.0723	0.73	41%	18.4	1.22	2.18	4.32	1.06	2.23
8. 7.1Final Survivor	thirsty_stallman	2026-10-04 16:10	0.920	$0.0962	$0.0648	0.67	48%	17.2	0.94	1.57	3.76	0.94	1.95
9. 7.1Final Survivor	sweet_dijkstra	2026-10-04 16:25	0.940	$0.1253	$0.0883	0.70	43%	20.6	1.14	2.18	4.36	1.00	2.34

Run times from: file time (finish).

Warnings

- 6.1 survivor: list cost per question ranges from $0.0358 to $0.0521 across its own runs (46% apart).

- 7.1Final Survivor: list cost per question ranges from $0.0962 to $0.1253 across its own runs (30% apart).

- Cached share ranges from 16% to 49% across runs (it grows with repeats of the same questions). Billed cost depends on run order; compare list cost.

- The decomposer's median call took 0.94-1.32 s across runs (41% apart). Its prompt is the same in every configuration, so time differences of this size are the provider's speed, not the system.

- Runs without telemetry (quality only; cost and time from the remaining runs): 5.1 survivor.

7. Safeguard markers

Configuration	Question-runs	Evidence notice	Noticed: pass / partial / fail	Grounding checked	Flagged	Flagged: pass / partial / fail	Check passed, judge failed	Input sanitized
6.1 survivor	75	9 (12%)	2/5/2	not checked	-	-	-	0
5.1 survivor	75	0 (0%)	0/0/0	not checked	-	-	-	0
7.1Final Survivor	75	18 (24%)	10/6/2	75	8 (11%)	5/2/1	1	0

A flagged answer the judge passed is a false alarm the user still sees (policy 'surface' labels it). 'Check passed, judge failed' counts wrong answers the deterministic check cannot see.

Spans the grounding check could not find in the evidence:

Span	Kind	Times	Judge: pass / partial / fail	Configurations
Command Module Columbia	name	3	3/0/0	7.1Final Survivor
A computed cross-ceremony intersection of nominee people across all three ceremonies is not supplied	quotation	1	0/1/0	7.1Final Survivor
Academy Awards.html	name	1	0/0/1	7.1Final Survivor
Estimates vary from 316,000 Haitian government to 222,570 UN OCHA estimate to 158,000 Medicine, Conflict and Survival to between 46,000 and 85,000 report commissioned by USAID	quotation	1	1/0/0	7.1Final Survivor
There is no computed result stating, for each intersecting person, whether they won at one or more of the three ceremonies, with the corresponding award and work	quotation	1	0/1/0	7.1Final Survivor
There is no deterministic result listing every person present at all three ceremonies with at least one nominated work from each ceremony	quotation	1	0/1/0	7.1Final Survivor
a computed intersection of nominee persons across the 68th, 70th, and 74th Academy Awards is not provided	quotation	1	0/1/0	7.1Final Survivor
at least one fatality estimate shown exceeds 100,000	quotation	1	1/0/0	7.1Final Survivor
could not complete	quotation	1	0/1/0	7.1Final Survivor
dated 1976 or later	quotation	1	1/0/0	7.1Final Survivor
disproving the aether	quotation	1	1/0/0	7.1Final Survivor
every person present at all three ceremonies	quotation	1	0/1/0	7.1Final Survivor

8. Routing

Not requested (--route-cheap and --route-strong).

Configuration	Route	Question-runs (share)	Quality	List $/q	Seconds/q
7.1Final Survivor	cheap	36 (48%)	1.000	$0.0325	10.0
7.1Final Survivor	strong	39 (52%)	0.872	$0.1753	26.8